

COMP9319 Web Data Compression and Search

Distributed path queries,
Compressed inverted index

1

Intro to distributed query evaluation

Web data is inherently distributed

Reuse some techniques from distributed
RDBMS if some schema info is known

**New techniques required if no schema info
is known**


In XML, these links are denoted in XLinks and
XPointers.

Example query

Assume data are distributed in 3 sites

Assume the RPE: $a.b^*.c$

Assume the query starts

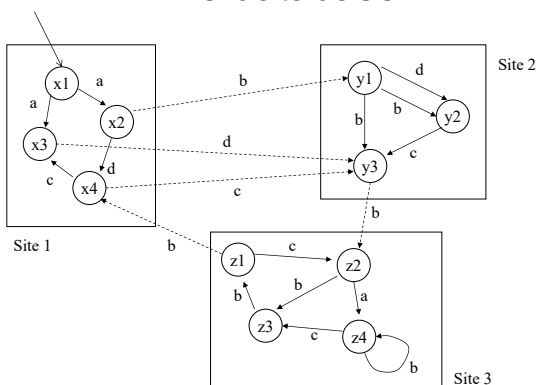


Regular path expressions

Regular expressions for path, e.g.:

Assignment Project Exam Help
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

The database



Naïve approach

A naïve approach takes too many
communication steps

=> we have to do more work locally

A better approach needs to

1. identify all external references
2. identify targets of external references

Input and output nodes

Site 1

Inputs: x1 (root), x4

Outputs: y1, y3

Site 2

Inputs: y1, y3

Outputs: z2

Site 3

Inputs: z2

Outputs: x4

Query Processing

Given a query, we compute its automaton

Send it to each site


Start an identical process at each site

Compute two sets $\text{Stop}(n, s)$ and $\text{Result}(n, s)$

Transmits the relations to a central location
and get their union

Stop and Result at site 2

Start	Stop
(y1, s2)	(z2, s2)
(y3, s2)	(z2, s2)



Start	Result
(y1, s2)	y2
(y1, s3)	y1
(y3, s3)	y3

Union the relations from all sites

Start	Stop
(y1, s2)	(z2, s2)
(y3, s2)	(z2, s2)

Start	Result
(x1, s3)	x1
(y4, s1)	x3
(y4, s1)	x4
(y1, s2)	y3
	y1
	y2
(z2, s2)	z2
(z2, s3)	z2

The result of the query
is {y3, z2, x3}

COMP9319: Web Data Compression and Search

Inverted index revisit & its compression

Slides modified from Hinrich Schütze and Christina Lioma slides on IIR

Inverted Index

For each term t , we store a list of all documents that contain t .

BRUTUS	→	1	2	4	11	31	45	173	174
CAESAR	→	1	2	4	5	6	16	57	132 ...
CALPURNIA	→	2	31	54	101				

⏟
dictionary

⏟
postings

Inverted index construction

- 1 Collect the documents to be indexed:

Friends, Romans, countrymen.

So let it be with Caesar

 ...
- 2 Tokenize the text, turning each document into a list of tokens:

Friends

Romans

countrymen

So

 ...
- 3 Do linguistic preprocessing, producing a list of normalized tokens, which are the indexing terms:

friend

roman

countryman

so

 ...
- 4 Index the documents that each term occurs in by creating an inverted index, consisting of a dictionary and postings.

4

Tokenizing and preprocessing

Doc 1. I did enact Julius Caesar: I was killed i' the Capitol; Brutus killed me.

Doc 2. So let it be with Caesar. The noble Brutus hath told you Caesar was ambitious:



Doc 1. i did enact julius caesar i was killed i' the capitol brutus killed me

Doc 2. so let it be with caesar the noble brutus hath told you caesar was ambitious

5

Generate postings

Doc 1. i did enact julius caesar i was killed i' the capitol brutus killed me
Doc 2. so let it be with caesar the noble brutus hath told you caesar was ambitious

term	docID
i	1
did	1
enact	1
julius	1
caesar	1
i	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2

6

Sort postings

7

Assignment Project Exam Help

<https://eduassistpro.github.io/>
 Add WeChat edu_assist_pro

Create postings lists, determine document frequency

term	docID	term	doc. freq.	postings lists
ambitious	2	ambitious	1	[2]
be	2	be	1	[2]
brutus	1	brutus	2	[1, 2]
brutus	2	brutus	2	[1, 2]
capitol	1	capitol	1	[1]
caesar	1	caesar	2	[1, 2]
caesar	2	caesar	2	[1, 2]
caesar	2	caesar	2	[1, 2]
did	1	did	1	[1]
enact	1	enact	1	[1]
hath	1	hath	1	[1]
i	1	i	1	[1]
i	1	i	1	[1]
i'	1	i'	1	[1]
it	2	it	1	[2]
julius	1	julius	1	[1]
killed	1	killed	1	[1]
killed	1	killed	1	[1]
let	2	let	1	[2]
me	1	me	1	[1]
noble	1	noble	1	[1]
so	2	so	1	[2]
the	1	the	2	[1, 2]
the	2	the	2	[1, 2]
told	2	told	1	[2]
you	2	you	1	[2]
was	1	was	2	[1, 2]
was	2	was	2	[1, 2]
with	2	with	1	[2]

8

Split the result into dictionary and postings file

BRUTUS	→	1	2	4	11	31	45	173	174	
CAESAR	→	1	2	4	5	6	16	57	132	...
CALPURNIA	→	2	31	54	101					
⋮										

dictionary

postings


9

Simple conjunctive query (two terms)

- Consider the query: BRUTUS AND CALPURNIA
- To find all matching documents using inverted index:
 - 1 Locate BRUTUS in the dictionary
 - 2 Retrieve its postings list from the postings file
 - 3 Locate CALPURNIA in the dictionary
 - 4 Retrieve its postings list from the postings file
 - 5 Intersect the two postings lists
 - 6 Return intersection to user

10

Intersecting two posting lists



- This is linear in the length of the postings lists.
- Note: This only works if postings lists are sorted.

11

Intersecting two posting lists

Typical query optimization

- Example query: BRUTUS AND CALPURNIA AND CAESAR
- Optimization: Process in order of increasing length of postings list, then keep cutting further
- For CAESAR, then CALPURNIA, then BRUTUS

12


13

Optimized intersection algorithm for conjunctive queries

```
INTERSECT( $\langle t_1, \dots, t_n \rangle$ )
1  terms ← SORTBYINCREASINGFREQUENCY( $\langle t_1, \dots, t_n \rangle$ )
2  result ← postings(first(terms))
3  terms ← rest(terms)
4  while terms ≠ NIL and result ≠ NIL
5  do result ← INTERSECT(result, postings(first(terms)))
6  terms ← rest(terms)
7  return result
```

14

Recall basic intersection algorithm



- Linear in the length of the postings lists.
- Can we do better?


15

Skip pointers

- Skip pointers allow us to **skip** postings that will not figure in the search results.
- This makes intersecting postings lists more efficient.
- Some postings lists contain several million entries – so efficiency can be an issue even if basic intersection is linear.
- Where do we put skip pointers?
- How do we make sure intersection results are correct?

16

Basic idea



17

Skip lists: Larger example

Intersection with skip pointers

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

18

19

Where do we place skips?

- Tradeoff: number of items skipped vs. frequency skip can be taken
- More skips: Each skip pointer skips only a few items, but we can frequently use it.
- Fewer skips: Each skip pointer skips many items, but we can not use it very often.

20

Phrase queries

- We want to answer a query such as [stanford university] – as a phrase.
- Thus *The inventor Stanford Ovshinsky never went to university* should **not** be a match.
- The concept of phrase query has proven easily understood by users.
- About 10% of web queries are phrase queries.
- Consequence for inverted index: it no longer suffices to store docIDs in postings lists.
- Two ways of extending the inverted index:
 - biword index (cf. COMP6714)
 - positional index

21

Positional indexes

- Postings lists in a **nonpositional** index: each posting is just a docID
- Postings lists in a **positional** index: each posting is a docID and a list of positions

22

Positional indexes: Example

Query: "*to₁ be₂ or₃ not₄ to₅ be₆*"

TO, 993427:

1: <7, 18, 33, 72, 86, 231>;

2: <1, 17, 74, 222, 255>;

4: <8, 16, 190, 429, 433>;

5: <363, 367>;

7: <13, 23, 191>; . . .

BE, 178239:

1: <17, 25>;

4: <17, 191, 291, 430, 434>;

5: <14, 19, 101>; . . . Document 4 is a match!

23

Inverted index

Dictionaries

data structure for storing the term

data

- Dictionary: the data structure for storing the term vocabulary

25

Dictionary as array of fixed-width entries

- For each term, we need to store a couple of items:
 - document frequency
 - pointer to postings list
 - . . .
- Assume for the time being that we can store this information in a fixed-length entry.
- Assume that we store these entries in an array.

26

Dictionary as array of fixed-width entries

term	document frequency	pointer to postings list
a	656,265	→
aachen	65	→
...
zulu	221	→

space needed: 20 bytes 4 bytes 4 bytes

How do we look up a query term q_i in this array at query time?
That is: which data structure do we use to locate the entry (row) in the array where q_i is stored?

27

Data structures for looking up term

- Two main classes of data structures: hashes and trees
- Some IR systems use hashes, some use trees.
- Criteria for when to use hashes vs. trees:
 - Is there a fixed number of terms or will it keep growing?
 - What are the relative frequencies with which various keys will be accessed?
 - How many terms are we likely to have?

28

Hashes

- Each vocabulary term is hashed into an integer.
- Try to avoid collisions
- At query time, do the following: hash query term, resolve collisions, locate entry in fixed-width array
- Pros: Lookup in a hash is faster than lookup in a tree.
 - Lookup time is constant.
- Cons
 - no way to find minor variants (*resume* vs. *résumé*)
 - no prefix search (all terms starting with *automat*)
 - need to rehash everything periodically if vocabulary keeps growing

29

Trees

- Trees solve the prefix problem (find all terms starting with *automat*).
- Simplest tree: binary tree
- Search is slightly slower than in the size of the vocabulary.
- $O(\log M)$ only holds for **balanced** trees.
- Rebalancing binary trees is expensive.
- B-trees** mitigate the rebalancing problem.
- B-tree definition: every internal node has a number of children in the interval $[a, b]$ where a, b are appropriate positive integers, e.g., $[2, 4]$.

30

Sort-based index construction

- As we build index, we parse docs one at a time.
- The final postings for any term are incomplete until the end. We must store all postings in memory and then do the sort in order to produce the final postings entry, we need a lot of space for the sort. This construction does not scale for large vocabularies or intermediate results on disk.

31

Same algorithm for disk?

- Can we use the same index construction algorithm for larger collections, but by using disk instead of memory?
- No: Sorting for example 100,000,000 records on disk is too slow – too many disk seeks.
- We need an **external** sorting algorithm.


32

“External” sorting algorithm (using few disk seeks)

- We must sort 100,000,000 non-positional postings.
 - Each posting has size 12 bytes (4+4+4: termID, docID, document frequency).
- Define a **block** to consist of 10,000,000 such postings
 - We can easily fit that many postings into memory.
 - We will have 10 such blocks.
- Basic idea of algorithm:
 - For each block: (i) accumulate postings, (ii) sort in memory, (iii) write to disk
 - Then merge the blocks into one long sorted order.

33

Merging two blocks



34

Why compression in information retrieval?

- First, we will consider space for dictionary
 - Main motivation for dictionary compression: make it small enough to keep in main memory
- Then for the postings file
 - Motivation: reduce disk space needed, decrease time needed to read from disk
 - Note: Large search engines keep significant part of postings in memory
- We will devise various compression schemes for dictionary and postings.

35

Dictionary compression

- The dictionary is small compar
- But we want to keep it in mem
- Also: competition with other a
- onboard computers, fast startup time
- So compressing the dictionary is important.

36

Recall: Dictionary as an array of fixed-width entries

vtas 4 bytes
= 11.2 MB


37

Fixed-width entries are bad.

- Most of the bytes in the term column are wasted.
 - We allot 20 bytes for terms of length 1.
- We can't handle HYDROCHLOROFLUOROCARBONS and SUPERCALIFRAGILISTICEXPIALIDOCIOUS
- Average length of a term in English: 8 characters
- How can we use on average 8 characters per term?

38

Dictionary as a string



39

Space for dictionary as a string

- 4 bytes per term for frequency
- 4 bytes per term for pointer to postings list
- 8 bytes (on average) for term in string
- 3 bytes per pointer into string (need $\log_2 8 \cdot 400000 < 24$ bits to resolve $8 \cdot 400,000$ positions)
- Space: $400,000 \times (4 + 4 + 3 + 8) = 7.6\text{MB}$ (compared to 11.2 MB for fixed-width array)

40

Dictionary as a string with blocking

...7systile9syzygetic8syzygial6syzygy1tszaibelyite6szecin...

freq.	postings ptr.	term ptr.
9	→	
92	→	
5	→	
71	→	
12	→	
...

41

Space for dictionary as a string with blocking

- Example block size $k = 4$
- Where we used 4×3 bytes for blocking ...
- ... we now use 3 bytes for one indicating the length of each term.
- We save $12 - (3 + 4) = 5$ bytes per block.
- Total savings: $400,000/4 \times 5 = 0.5\text{MB}$
- This reduces the size of the dictionary from 7.6 MB to 7.1 MB.

42

Look up of a term with out blocking


Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

43

Lookup of a term with blocking: (slightly) slower



44

Front coding

One block in blocked compression ($k = 4$) ...
 8 automata8 automate9 automatic10 automation
 ↓
 ... further compressed with front coding.
 8 automa*t a 1 o e 2 o i c 3 o i n

45

Dictionary compression for Reuters: Summary

data structure	size in MB
dictionary, fixed-width	11.2
dictionary, term pointers into string	7.6
~, with blocking, $k = 4$	7.1
~, with blocking & front coding	5.9

46

Postings compression

- The postings file is much larger than the dictionary, factor of at least 10.
- Key desideratum: store each posting compactly
- A posting for our purposes is a docID.
- For Reuters (800,000 documents), we would use 32 bits per docID when using 4-byte integers.
- Alternatively, we can use $\log_2 800,000 \approx 19.6 < 20$ bits per docID.
- Our goal: use a lot less than 20 bits per docID.

47

Key idea: Store gaps instead of docIDs

- Each postings list is ordered in increasing order of docID.
- Example postings list: COMPUTER
- It suffices to store **gaps**: 283159-
- Example postings list using gaps
- Gaps for frequent terms are small.
- Thus: We can encode small gaps with fewer than 20 bits.

48

Gap encoding

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

49

Variable length encoding

- Aim:
 - For ARACHNOCENTRIC and other rare terms, we will use about 20 bits per gap (= posting).
 - For THE and other very frequent terms, we will use only a few bits per gap (= posting).
- In order to implement this, we need to devise some form of **variable length encoding**.
- Variable length encoding uses few bits for small gaps and many bits for large gaps.

50

Variable byte (VB) code

- Used by many commercial/research systems
- Good low-tech blend of variable-length coding and sensitivity to alignment matches (bit-level codes, see later).
- Dedicate 1 bit (high bit) to be a **continuation bit** c .
- If the gap G fits within 7 bits, binary-encode it in the 7 available bits and set $c = 1$.
- Else: encode lower-order 7 bits and then use one or more additional bytes to encode the higher order bits using the same algorithm.
- At the end set the continuation bit of the last byte to 1 ($c = 1$) and of the other bytes to 0 ($c = 0$).

51

Properties of gamma code

- Gamma code is prefix-free
- The length of offset is $\lfloor \log_2 G \rfloor$ bits.
- The length of length is $\lfloor \log_2 G \rfloor + 1$ bits,
- So the length of the entire code is $2 \times \lfloor \log_2 G \rfloor + 1$ bits.
- γ codes are always of odd length.
- Gamma codes are within a factor of 2 of the optimal encoding length $\log_2 G$.

58

Gamma codes: Alignment

- Machines have word boundaries – 8, 16, 32 bits
- Compressing and manipulating at granularity of bits can be slow.
- Variable byte encoding is aligned and thus potentially more efficient.
- Regardless of efficiency, variable byte is conceptually simpler at little additional space cost.

59

Compression of Reuters

data structure	size in MB
dictionary, fixed-width	
dictionary, term pointers into stri	
~, with blocking, k = 4	
~, with blocking & front coding	
collection (text, xml markup etc)	
collection (text)	960.0
T/D incidence matrix	40,000.0
postings, uncompressed (32-bit word)	400.0
postings, uncompressed (20 bits)	250.0
postings, variable byte encoded	116.0
postings, gamma encoded	101.0

60

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro